

Adaptation Experiment 1: Univariate Control

1 Theoretical overview and goal

A forecasting adapter is useful only if it improves a well-defined, untouched backbone. For a look-back $x \in \mathbb{R}^L$, a pretrained forecaster produces $\hat{y}_0 = f(x) \in \mathbb{R}^H$. The univariate control evaluates f directly, without retrieved neighbors, gates, or trainable adaptation. It therefore separates backbone error from error introduced by the retrieval and adaptation pipeline.

The main metrics are

$$\text{MSE} = H^{-1} \|\hat{y}_0 - y\|_2^2, \quad \text{nMSE} = H^{-1} \left\| \frac{\hat{y}_0 - y}{\sigma(x) + \varepsilon} \right\|_2^2.$$

Per-user aggregation and the worst 10% user mean reveal whether a good global mean hides systematic failures on a subset of series.

Experimental goal. Establish the no-adaptation reference, measure its temporal and user-wise heterogeneity, and provide the denominator against which every retrieval method must be compared.

2 Implementation details

Dates are split chronologically as $(T_0, T_1, T_2, T_3) = (0.30, 0.35, 0.15, 0.20)$. The control command maps T_1 , T_2 , and T_3 to the labels **train**, **oracle**, and **eval**; the Slurm job evaluates the held-out T_3 period. Each query contains all users at one date, but the backbone is called as a univariate model with one channel per sample.

Dataset preparation automatically reads the sibling **config.json**. Portable top-level options and adaptation-scoped options are combined before the chronological split; user exclusions are additive and the selected config path is logged. The same preparation contract is used by extraction, the univariate control, and all downstream adapters.

The test profile uses Electricity 168–24 with sparse evaluation. The small profile below is the current Chronos pilot; the large profile retains Chronos but widens the control to every dataset and setting used by Adaptation.

Component	Current sweep value
Datasets	Electricity, Solar
Look-back-horizon	168–24, 672–168
Backbone	Chronos
Input normalization	Instance normalization
Evaluation stride	128 dates
Seed	1

Implementation: `src/experiments/experiment_univariate.py`

Launcher: `univariate.slurm`

It writes window-level losses, a JSON summary, a tensor payload, error distributions, horizon-error curves, user scatter plots, and prediction examples. The primary comparison is the same T_3 MSE/nMSE later used by adapted methods.

3 Interpretation

This experiment is descriptive rather than trainable. Any adapted method must beat the control on the same queries and aggregation rule. Improvements on T_1 or T_2 alone are insufficient evidence because those periods are used to fit adapters or gates.